

Understanding Numba

Valentin Haenel

Senior Software Engineer - Open Source

<https://anaconda.com>

LLVM Social Berlin - Wednesday 23 October 2024

Version: 2024-10-23-llvm-berlin-meetup-3-g542a4f87c7

<https://github.com/esc/numba-talk>



This work is licensed under the *Creative Commons Attribution-ShareAlike 3.0 License*.

- Senior Software Engineer at Anaconda
- Working on Numba full-time
- Doing this for over 5 years
- Actively involved in:
 - Typed containers
 - Release and community management
 - Compiler frontend design and implementation

Outline

- 1 Introduction
- 2 Going Deeper
- 3 `llvmlite`
- 4 NumPy Support
- 5 Tips and Tricks
- 6 Compiler Toolkit
- 7 Summary



Numba in a Nutshell

- A compiler that might make your code faster
- Requires importing a decorator: `@jit`
- And decorating functions with it
- Numba = NumPy + Mamba (fast snake)

Numba Explained

- Numba is a
 - just-in-time
 - type-specializing
 - function compiler
 - for accelerating numerically-focused Python

- LLVM is a compiler toolkit
- Numba uses it as a compiler backend
- Access via `llvmlite`

Example

- Sieve of Eratosthenes

	2	3	4	5	6	7	8	9	10	Prime numbers			
11	12	13	14	15	16	17	18	19	20	2	3	5	7
21	22	23	24	25	26	27	28	29	30	11	13	17	19
31	32	33	34	35	36	37	38	39	40	23	29	31	37
41	42	43	44	45	46	47	48	49	50	41	43	47	53
51	52	53	54	55	56	57	58	59	60	59	61	67	71
61	62	63	64	65	66	67	68	69	70	73	79	83	89
71	72	73	74	75	76	77	78	79	80	97	101	103	107
81	82	83	84	85	86	87	88	89	90	109	113		
91	92	93	94	95	96	97	98	99	100				
101	102	103	104	105	106	107	108	109	110				
111	112	113	114	115	116	117	118	119	120				

Example

```
import numpy as np
from numba import jit

@jit # simply add the jit decorator
def primes(max=100000):
    numbers = np.ones(max, dtype=np.uint8) # initialize the boolean sieve
    for i in range(2, max):
        if numbers[i] == 0: # has previously been crossed off
            continue
        else: # it is a prime, cross off all multiples
            x = i + i
            while x < max:
                numbers[x] = 0
                x += i
    # return all primes, as indicated by boolean positions that have value 1
    return np.nonzero(numbers)[0][2:]
```

Example

```
In [5]: %timeit sieve.primes.py_func()
124 ms  $\pm$  2.72 ms per loop (mean  $\pm$  std. dev. of 7 runs, 10 loops each)

In [6]: %timeit sieve.primes()
308  $\mu$ s  $\pm$  8.93  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs, 1000 loops each)
```

Open Source Status

- Several companies and many open source contributors
 - 5 FTE funded to contribute to Numba between Anaconda, Intel, Quansight, Nvidia and others
 - 7-12 non-core contributors per release
- Issue and PR lists growing
- Very active community on GitHub, Discourse and Matrix/Gitter

Community Usage

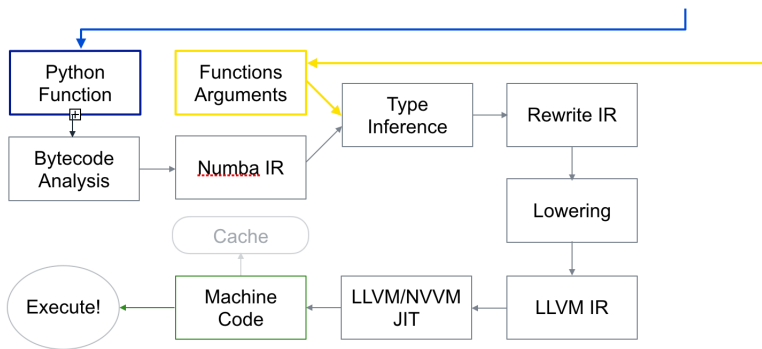
- 10s of Millions of Downloads per month
- On GitHub about 120k repositories are listed as dependents (Oct 2024)
- Several high-profile libraries use it:
 - PyData Sparse → sparse matrix implementation
 - UMAP → Uniform Manifold Approximation
 - Tardis → Super Nova Simulator
 - Pandas → Dataframe Library
 - Open AI Whisper → Speech-to-Text
 - NVIDIA RAPIDS → GPU Library
 - many, many more, we eventually stopped keeping track

Outline

- 1 Introduction
- 2 Going Deeper
- 3 `llvmlite`
- 4 NumPy Support
- 5 Tips and Tricks
- 6 Compiler Toolkit
- 7 Summary

Numba Flow

```
@jit
def do_math(a, b):
    ...
>>> do_math(x, y)
```



- Translate Python objects of supported types into representations with no CPython dependencies ("unboxing")
- Compile Python bytecode from your decorated function into machine code.
- Swap calls to builtins and NumPy functions for implementations provided by Numba (or 3rd party Numba extensions)
- Allow LLVM to inline functions, autovectorize loops, and do other optimizations you would expect from a C compiler
- Allow LLVM to exploit the instruction sets of your hardware (SSE, AVX)
- When calling the function, release the GIL if requested
- Convert return values back to Python objects ("boxing")

What Numba does not do

- Automated translation of CPython or NumPy implementations
- Automatic compilation of 3rd party libraries
- Partial compilation
- Automatic conversion of arbitrary Python types
- Change the layout of data allocated in the interpreter
- Translate entire programs
- Magically make individual NumPy functions faster

When is Numba unlikely to help?

- Whole program compilation
- Critical functions have already been converted to C or optimized Cython
- Need to interface directly with C++
- Algorithms that are not primarily numerical, e.g. string manipulation
- Exception: Numba can do pretty well at bit manipulation

Outline

- 1 Introduction
- 2 Going Deeper
- 3 `llvmlite`
- 4 NumPy Support
- 5 Tips and Tricks
- 6 Compiler Toolkit
- 7 Summary

- Numba's LLVM binding package
- A "lightweight" Python wrapper
- Not exactly intended as a standalone package, but a few people use it anyway
- Separation of concerns, easier to develop this way
- Uses the C++ API via C – easier to bind to from Python
- Stability is key, we only upgrade LLVM if necessary, don't chase the bleeding edge

llvmlite releases

- Released in lock-step with Numba
- About two to four times a year
- Each release only supports one (sometimes two) LLVM versions
 - → several disgruntled users
- Currently we support 14 but transitioning to 15
- ZERO-based versioning
- Patch releases only to fix severe regressions
- No backports

llvmlite packages

- Numba team creates two types of packages, wheels and conda packages
- The Numba team statically distributes LLVM
- Python wheels for distribution via PyPI
 - 20 * 30-40 MB
- Conda packages for distribution via the numba channel on anaconda.org
 - 25 * 20-40 MB
- Other distributors choose to link dynamically
- (This often creates issues for distributors, since they also only want to support one LLVM version, usually not the one we support but the latest)
- Hot take: LLVM is not a shared library – but something to be vendored

Why not llvmpy

- llvmlite comes from an era of LLVM 3.something
- Stability was more important, Numba is the main consumer
- Only the JIT part of LLVM needed (MCJIT)
- For example: our builder-API is entirely in Python and string based
- Better to control the bindings ourselves

Builder-API Example

```
from llvmlite import ir

# Create some useful types
double = ir.DoubleType()
fnty = ir.FunctionType(double, (double, double))

# Create an empty module...
module = ir.Module(name=__file__)
# and declare a function named "fpadd" inside it
func = ir.Function(module, fnty, name="fpadd")

# Now implement the function
block = func.append_basic_block(name="entry")
builder = ir.IRBuilder(block)
a, b = func.args
result = builder.fadd(a, b, name="res")
builder.ret(result)

print(module)
```

Builder-API Example

```
; ModuleID = "examples/ir_fpadd.py"
target triple = "unknown-unknown-unknown"
target datalayout = ""

define double @"fpadd"(double %".1", double %".2")
{
entry:
    %"res" = fadd double %".1", %".2"
    ret double %"res"
}
```


Code Generation Example

```
from __future__ import print_function

from ctypes import CFUNCTYPE, c_double

import llvmlite.binding as llvm

# All these initializations are required for code generation!
llvm.initialize()
llvm.initialize_native_target()
llvm.initialize_native_asmprinter() # yes, even this one
```

Code Generation Example

```
llvm_ir = """
; ModuleID = "examples/ir_fpadd.py"
target triple = "unknown-unknown-unknown"
target datalayout = ""

define double @"fpadd"(double %".1", double %".2")
{
entry:
    %"res" = fadd double %".1", %".2"
    ret double %"res"
}
"""
```

Code Generation Example

```
def create_execution_engine():  
    """  
    Create an ExecutionEngine suitable for JIT code generation on  
    the host CPU. The engine is reusable for an arbitrary number of  
    modules.  
    """  
    # Create a target machine representing the host  
    target = llvm.Target.from_default_triple()  
    target_machine = target.create_target_machine()  
    # And an execution engine with an empty backing module  
    backing_mod = llvm.parse_assembly("")  
    engine = llvm.create_mcjit_compiler(backing_mod, target_machine)  
    return engine
```

Code Generation Example

```
def compile_ir(engine, llvm_ir):  
    """  
    Compile the LLVM IR string with the given engine.  
    The compiled module object is returned.  
    """  
    # Create a LLVM module object from the IR  
    mod = llvm.parse_assembly(llvm_ir)  
    mod.verify()  
    # Now add the module and make sure it is ready for execution  
    engine.add_module(mod)  
    engine.finalize_object()  
    engine.run_static_constructors()  
    return mod
```

Code Generation Example

```
engine = create_execution_engine()
mod = compile_ir(engine, llvm_ir)

# Look up the function pointer (a Python int)
func_ptr = engine.get_function_address("fpadd")

# Run the function via ctypes
cfunc = CFUNCTYPE(c_double, c_double, c_double)(func_ptr)
res = cfunc(1.0, 3.5)
print("fpadd(...) =", res)
```

Outline

- 1 Introduction
- 2 Going Deeper
- 3 llvmlite
- 4 NumPy Support**
- 5 Tips and Tricks
- 6 Compiler Toolkit
- 7 Summary

Implementing NumPy

- Numba does not use much of the NumPy C implementations
- (Only some ufuncs are "borrowed".)
- We implement the NumPy API using Numba compatible/supported Python
- Treat NumPy as DSL for array oriented computing

Implementing Numpy

- Implement: `numpy.linalg.norm`
- For vectors, `ord` is:
 - `inf` \rightarrow `min(abs(x))`
 - `0` \rightarrow `sum(x != 0)`
- Implemented in: `numba.targets.linalg.py`

Implementing Numpy

```
elif ord == -np.inf:
    # min(abs(a))
    ret = abs(a[0])
    for k in range(1, n):
        val = abs(a[k])
        if val < ret:
            ret = val
    return ret

elif ord == 0:
    # sum(a != 0)
    ret = 0.0
    for k in range(n):
        if a[k] != 0.:
            ret += 1.
    return ret
```

Outline

- 1 Introduction
- 2 Going Deeper
- 3 llvmlite
- 4 NumPy Support
- 5 Tips and Tricks**
- 6 Compiler Toolkit
- 7 Summary

Tips and Tricks

- Always use `@jit`
- Prefer NumPy arrays for numerical data
- Use typed containers from `numba.typed` for nested data
- `for` loops are fine
- Array expressions are fused if on the same line, so those are fine too

Typed Containers

```
from numba import njit
from numba.typed import List

l = List()  # instantiate a new typed list
l1 = List(); [l1.append(i) for i in range(5)]
l.append(l1)  # add the first sub-list
l2 = List(); [l2.append(i) for i in range(10)]
l.append(l2)  # add the second sub-list
print(l)

@njit
def func(my_list):
    # modify list in a compiled context
    for i in range(10):
        my_list[1][i] = 23

func(l)
print(l)
```

Typed Containers

```
$ python code/typed.py
```

```
[[0, 1, 2, 3, 4], [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]]
```

```
[[0, 1, 2, 3, 4], [23, 23, 23, 23, 23, 23, 23, 23, 23, 23]]
```

Fused Expressions

```
import numpy as np
from numba import njit

a, b = np.arange(1e6), np.arange(1e6)

@njit
def func(a, b):
    return a*b-4.1*a > 2.5*b
```

Fused Expressions

```
In [1]: from fused import a,b,func
```

```
In [2]: %timeit func.py_func(a,b)
```

```
4.68 ms ± 89.7 µs per loop (mean ± std. dev. of 7 runs, 100 loops each)
```

```
In [3]: %timeit func(a,b)
```

```
626 µs ± 22.4 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)
```

OS and Hardware Support

- Windows 7 and later 64-bit only
- macOS 10.9 and later, x86_64 (Intel) and arm64 (M1, M2 etc...)
- Linux x86_64 and aarch64 (no ppc64le anymore)
- NVIDIA CUDA GPUs (Compute capability 5.3 and later, CUDA 11.2 and later)
 - → currently being refactored into separate package numba-cuda

Python versions

- Python 3.10 - 3.13 as of Numba 0.62 and llvmlite 0.43

Packaging

- You can depend on Numba to perform the heavy lifting!
- We run CI on most Python/NumPy/OS/Hardware combinations
- You can ship a single source package
 - PyPi
 - anaconda.org
- No need to pre-compile binaries for your users

Outline

- 1 Introduction
- 2 Going Deeper
- 3 llvmlite
- 4 NumPy Support
- 5 Tips and Tricks
- 6 Compiler Toolkit**
- 7 Summary

inspect methods

- There are various inspect methods on compiled functions
 - `inspect_types` → Prints out results from type-inference
 - `inspect_llvm` → Obtain LLVM IR
 - `inspect_asm` → Obtain print out of function assembly
 - `inspect_cfg` → Obtain Control Flow Graph
 - `inspect_dissam_cfg` → Control Flow Graph and disassembly from reversing generated ELF

Custom Compiler Passes

```
from numba import njit
from numba.core import ir
from numba.core.compiler import CompilerBase, DefaultPassBuilder
from numba.core.compiler_machinery import FunctionPass, register_pass
from numba.core.untyped_passes import IRProcessing
from numbers import Number

# Register this pass with the compiler framework, declare that it will not
# mutate the control flow graph and that it is not an analysis_only pass (it
# potentially mutates the IR).
@register_pass(mutates_CFG=False, analysis_only=False)
class ConstsAddOne(FunctionPass):
    _name = "consts_add_one" # the common name for the pass

    def __init__(self):
        FunctionPass.__init__(self)
```

Custom Compiler Passes

```
# implement method to do the work, "state" is the internal compiler
# state from the CompilerBase instance.
def run_pass(self, state):
    func_ir = state.func_ir # get the FunctionIR object
    mutated = False # used to record whether this pass mutates the IR
    # walk the blocks
    for blk in func_ir.blocks.values():
        # find the assignment nodes in the block and walk them
        for assgn in blk.find_insts(ir.Assign):
            # if an assignment value is a ir.Consts
            if isinstance(assgn.value, ir.Const):
                const_val = assgn.value
                # if the value of the ir.Const is a Number
                if isinstance(const_val.value, Number):
                    # then add one!
                    const_val.value += 1
                    mutated |= True
    return mutated # return True if the IR was mutated, False if not.
```

Custom Compiler Passes

```
class MyCompiler(CompilerBase): # custom compiler extends from CompilerBase

    def define_pipelines(self):
        # define a new set of pipelines (just one in this case)
        pm = DefaultPassBuilder.define_nopython_pipeline(self.state)
        # Add the new pass to run after IRProcessing
        pm.add_pass_after(ConstsAddOne, IRProcessing)
        pm.finalize()
        return [pm]
```

Custom Compiler Passes

```
@njit(pipeline_class=MyCompiler) # JIT compile using the custom compiler
def foo(x):
    a = 10
    b = 20.2
    c = x + a + b
    return c

print(foo(100)) # 100 + 10 + 20.2 (+ 1 + 1), extra + 1 + 1 from the rewrite!
```


Outline

- 1 Introduction
- 2 Going Deeper
- 3 llvmlite
- 4 NumPy Support
- 5 Tips and Tricks
- 6 Compiler Toolkit
- 7 Summary**

Understood Numba?

- Numba is a
 - just-in-time
 - type-specializing
 - function compiler
 - for accelerating numerically-focused Python

What is Cooking?

- AST/Source based frontend
- RVSDG based lambda-calculus intermediary representation
- Hindley-Milner style type inference
- egglog/Equality saturation
- PIXIE based backend, toolchain and executable format
- Transition from MCJIT \rightarrow ORCJIT
- Compiler modularization: 'numba-cuda' and 'numba-numpy'
- Refactoring to support NumPy 2.0 type system

- Effectively: complete project overhaul

Getting in Touch

- <https://numba.pydata.org>
- <https://github.com/numba>
- <https://numba.discourse.group/>
- Preferred: GitHub + Gitter + Discourse
- Stuck? It can't hurt to ask. ;-)